

Redesigning intake gives each of the 20 claim handlers back 52 minutes a day

17.5 hours a day across the 20 of them — 2.8 full-time people · two thirds of that work goes to claims that were always going to close quietly · a person still decides anything adverse

TODAY

What one claim costs one claim handler

40 minutes per claim

- Re-reads the report, opens attachments
- Chases a missing CPR
- Switches EASY · IDB · e-mail · paper
- +30% for breaks and interruptions

REDESIGNED

The same checks and the same exits, reordered

19 minutes per claim

- 52% cut, not the 75% max
- 4 AI agents, 1 rules-only step
- 2 human checkpoints kept
- Handlers re-check before they trust it

WHERE THE EFFORT ACTUALLY GOES

- 5%** **Models and machines**
€0.08 a claim · €2,400 a year
- 25%** **Plumbing and audit trail**
EASY/IDB wiring, logging every decision
- 70%** **Process, training, trust**
adoption is the real schedule risk

Follows BCG's 10-20-70 rule, adjusted for this process

52 min

freed per claim handler, every working day (of 6.4 productive hours)

17.5 h

across the 20 claim handlers — 2.8 full-time people

€153k

value of that capacity, per year

14 mths

to pay back the one-off build of €146k

What “€153k” means: 4,370 handler-hours a year × €34.92, the loaded cost of one handler-hour (€45,100 gross plus employer contributions, over actual working hours). Capacity, not cash — it becomes money only if the hours go to the backlog or are absorbed by normal staff turnover.

So what: cost is not what decides this. The AI costs €2,400 a year to run; keeping it healthy takes a quarter of an engineer, €24,400 — so €26,800 a year all in, against €152,600 of capacity released. What decides it is governance: proving a person reviewed anything that can go against a claimant, and whether handlers trust the drafts.


Conservative throughout: the 21-minute-per-claim saving is a three-point estimate (best 30 / likely 21 / worst 12), not the design maximum. Build-up on slide 7, sources on slide 8.

Three ways in and a separate side-track diagram become one classified intake

FNOL as it runs today (left) against the redesign (right) · first report to an opened, policy-linked claim


Trigger / start Step / decision Exception / rework Exit

TODAY



AI agent (LLM call) Rules only, no model Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




These two agents together replace 3 entry channels, the auto-import failure path, and the entire 1.1 side-track diagram.

Three checks done one after another become one check done all at once

Is it complete, linked to the right policy, and not already in the system? None of the three depends on the others


- Trigger / start
- Step / decision
- Exception / rework
- Exit

TODAY



- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




The AI writes the draft; the handler still makes the decision

From a system warning to a closed claim or a fully opened one — nothing that can go against a claimant is automated

- Trigger / start
- Step / decision
- Exception / rework
- Exit


TODAY



Today: 7 decision gates, 3 entry channels, 10 exception or exit branches across the full FNOL flow.

- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




Net result across all three parts:


4 LLM agents + 1 rules-only node · 2 human checkpoints kept by design · 3 entry channels → 1 · 3 sequential gates → 1 parallel check
Coverage Check runs on the same state, beside claim creation — so the handoff between the two workflows disappears rather than being managed

Start where a mistake is cheap — that is what earns the trust the bigger wins depend on

“Hard” means build effort and regulatory exposure together, not just weeks of work



AES rulings and underwriting authority are deliberately absent — those decisions are not ours to make · frequencies from the case study




THE STACK



Deployed to the cloud GGW already use · infrastructure as code · one on/off switch per agent · model version pinned

HOW IT GOES LIVE



If an agent fails or is unsure, the claim drops into today's manual queue — slower, never wrong · what it costs is on slide 7

A · INTAKE TIME — MINUTES PER CLAIM, PER PERSON

Step	Today	Redesigned
Read the report, re-read it, open attachments	10.0	0.7
Take notes, check fields, chase a missing CPR	6.1	2.8
Attach policy · duplicate check · coverage date	5.0	0.0
Decide, draft and send / open the full claim	5.5	3.0
Work an escalation or an underwriting hold	1.0	0.6
Switch between EASY, IDB, e-mail and paper	3.2	0.7
Touch time	30.8	7.8
Breaks, fatigue, interruptions, rework (+30%)	9.2	2.3
Standard time — what the design allows	40.0	10.1
What the business case actually banks	40.0	19.0

The design allows 10.1 min a claim. We do not bank that. The case uses a three-point estimate — (best 30 + 4 × likely 21 + worst 12) ÷ 6 = 21.0 min saved a claim — which is the 19.0 min above, and the 19 minutes on slide 1. In plain terms: we assume handlers capture 70% of the saving the design allows.

Cross-check: 40 min × 12,500 claims a year = 8,332 handler-hours — 26% of the 20 handlers' year, 17% of all 30. Per handler: 2.5 claims a day at 40 min is 1.7 of a 6.4-hour day.

D · THE BUSINESS CASE
EVERY YEAR
Capacity released
Worth, at €34.92 per handler-hour
Cost to run it — AI models and Azure
Ongoing support — 0.25 of an engineer
Net value, every year

4,370 handler-h · 2.8 FTE
€152,600 a year
– €2,400
– €24,400
€125,800

ONE-OFF
Build — 3 people × 6 months, fully loaded
Payback

€146,400
14 months

The build is loaded team capacity, not new cash — the AI team is already funded. New cash needed: €2,400 a year. A shorter build costs less: 4 months ≈ €98k.

B · AI COST TO PROCESS ONE CLAIM

Agent	Tokens in	Out	Runs	US\$ / claim
Extraction	11,500	700	100%	0.0314
Classification	900	120	100%	0.0032
Validation — rules only, no model	—	—	100%	0.0000
Exception research	2,500	400	20%	0.0020
Outcome drafting	2,200	550	100%	0.0110
Subtotal				0.0476
× 1.75 for retries, guardrails, evals				0.0833
Per claim				≈ €0.077

One claim ≈ a scanned form, an e-mail and 3–4 attachments. All four agents run on one model tier — gpt-5.6 terra, \$2 in / \$12 out per million tokens.

C · AZURE, WEST EUROPE — MONTHLY

Component	US\$/mth
4 container apps — 2 always on, 2 burst only	55
PostgreSQL B2s + 64 GiB storage and backup	61
Blob storage for documents + audit logs	8
Net new, incl. ~8% West Europe premium	134

Registry, networking, Key Vault and DevOps (≈ €670/yr) already exist, so they are not charged here. Azure DevOps gives repos and pipelines, not a database — PostgreSQL is genuinely new.

E · WHO GETS THE TIME BACK — CONFIRMED BY GGW

20 claim handlers, 10 support staff (GGW, 8 Sep) · 52.4 min each × 20 = 17.5 h a day

Per claim handler	Today	Redesigned	Difference
Claims taken in, per day	2.5	2.5	—
Time on intake, per day	100 min	48 min	– 52 min
Share of a productive day	26%	12%	– 14 pts
Time back, each, per day			52 min

Who does intake

20 claim handlers process claims; 10 support staff besides

GGW (P. Gossmann), 8 Sep 2026

Failed EASY imports

Not noticed today — so detection comes before retries

GGW (P. Gossmann), 8 Sep 2026

Volumes and shares

~50 first reports a day · 65% minor · ~30 in the team

GGW case study brief

Wage base — handlers

Schadenregulierer/in, €45,100 gross — the role doing intake

StepStone Gehaltsreport

Employer on-cost

≈ €23 per €100 gross — health, pension, care

Lohnnebenkosten 2026, sevdesk

Working time, breaks

38 h tariff week; statutory break rules

Arbeitszeitgesetz §4

Build & upkeep pay

Senior ML engineer ≈ €98k; juniors at €70k as briefed

Glassdoor, ML Engineer Germany

Model prices

gpt-5.6 terra — \$2 in / \$12 out per million tokens

OpenAI API pricing

Container hosting

Consumption plan rates and the monthly free grant

Azure Container Apps pricing

Database

PostgreSQL Flexible Server B2s — 2 vCore / 4 GiB

Azure Database for PostgreSQL

Upkeep after go-live

15–20% of build effort a year — our 0.25 FTE is 17%

Software maintenance rule of thumb

Industry baseline

17 of 20 carriers still manual; none let AI decide alone

Decerto, 2026 Claims Decisioning Pulse

Where value comes from

The 10-20-70 split — algorithms, technology, people

BCG, AI in customer service operations

Fixing import failures

Retry vs. dead-letter, idempotency keys, alert on rate

Dead-letter queues and poison messages

STILL OUR ASSUMPTIONS — AND THE FIRST THING WE WOULD TEST

Every per-step minute on slide 7 · the 30% allowance for breaks and interruptions · how often a CPR is missing (30%) or underwriting is asked (25%) · how often the redesign raises an exception (20%) · the tokens a typical claim needs · the €146k build · the 0.25 FTE upkeep.

Who does intake is no longer among them — GGW answered it, and every per-person figure in this deck is rebuilt on that answer. Two weeks of watching live intake would replace the rest with measurements.

confirmed by GGW or stated in the brief

published source, linked

Prices checked September 2026 · US\$ converted at €0.92 · every euro figure is fully loaded — gross pay plus employer contributions