

Redesigning intake roughly doubles a claim handler's daily capacity

Each claim drops from 40 to 19 minutes. Same 20 handlers, same shift: that's a capacity ceiling, not confirmed extra demand.

TODAY

What one claim costs one claim handler

40 minutes per claim

- Re-reads the report, opens attachments
- Chases a missing CPR
- Switches EASY · IDB · e-mail · paper
- +30% for breaks and interruptions

REDESIGNED

Every check and every exit kept, run in a new order

19 minutes per claim

- 21 minutes back on every claim (a 52% cut)
- 4 AI agents, 1 rules-only step
- 2 human checkpoints kept
- Handlers re-check before they trust it

WHERE THE EFFORT ACTUALLY GOES

- 5%** **Models and machines**
€0.08 a claim · €2,400 a year
- 25%** **Plumbing and audit trail**
EASY/IDB wiring, logging every decision
- 70%** **Process, training, trust**
adoption is the real schedule risk

BCG's 10-20-70, tilted to plumbing: EASY and IDB are the constraint

52 min

freed per claim handler per day, out of 6.4 productive hours

17.5 h

across the 20 claim handlers, or 2.8 full-time people

€153k a year

what the 4,370 freed handler-hours are worth, at €34.92 fully loaded


14 mths

to pay back the one-off build of €146k

Recommendation: start with the failed imports nobody sees today, then intake. A person approves every decision that can go against a claimant, at every stage.

Seven stations share one claim record, so nothing is handed off or read twice

SYSTEM AT A GLANCE




AI drafts. Rules check. A person decides anything that can go against a claimant.

Three ways in and a separate side-track diagram become one classified intake

Today on the left, the redesign on the right


- Trigger / start
- Step / decision
- Exception / rework
- Exit

TODAY



- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE one shared claim record, one audit trail




These two agents replace 3 entry channels, the auto-import failure path, and the whole 1.1 side track.


Three checks done one after another become one check done all at once

None of the three checks depends on the other two

TODAY




AI-NATIVE one shared claim record, one audit trail



The AI writes the draft; the handler still makes the decision


Nothing that can go against a claimant is automated

TODAY



Today: 7 decision gates, 3 entry channels, 10 exception or exit branches across the full FNOL flow.


AI-NATIVE *one shared claim record, one audit trail*




Net result: 3 entry channels → 1 · 3 sequential gates → 1 parallel check · the Coverage Check handoff disappears, it runs on the same state

Start where a mistake is cheap. That is what earns the trust the bigger wins need.

“Hard” means build effort and regulatory exposure together, not just weeks of work




AES rulings and underwriting authority are deliberately absent. Those decisions are not ours to make.




THE STACK



OPERATING MODEL




HOW IT GOES LIVE



If an agent fails or is unsure, the claim falls back to today's manual queue. Slower, never wrong.

THE GRAPH



Every minute is observed, and the redesign pays for itself in 14 months

A · TIME PER CLAIM, PER HANDLER (MINUTES)

Step	Today	Redesigned
Read & extract the report	16.1	3.5
Check policy, CPR & escalations	6.0	0.6
Decide, draft and send	5.5	3.0
Switch between systems	3.2	0.7
Touch time	30.8	7.8
+30% breaks & interruptions	9.2	2.3
Standard time	40.0	10.1
What we bank	40.0	19.0

We bank 19 min, not the 10.1 the design allows: the business case runs on a three-point estimate of best, likely and worst case.

Basis: at 40 minutes a claim, today's standard time, an 8-hour shift holds capacity for 12 claims; at 19 (the banked, most-likely estimate), capacity is ~25, roughly double. That's a ceiling, not the 52 minutes freed above, which is actual time freed on today's claim volume. Confirmed by GGW (8 Sep 2026).

B · THE BUSINESS CASE

EVERY YEAR	
Capacity released	4,370 handler-h · 2.8 FTE
Worth, at €34.92 per handler-hour	€152,600
Cost to run it — AI models and Azure	– €2,400
Ongoing support — 0.25 of an engineer	– €24,400
Net value, every year	€125,800
ONE-OFF	
Build — 3 people × 6 months, fully loaded	€146,400
Payback	14 months

Running cost in full: ≈ €0.08/claim in AI tokens + ≈ \$134/month in Azure → €2,400/year, all in.

The build is loaded team capacity, not new cash. The AI team is already funded.

Every number is sourced, dated, or named as an assumption we'd test first

- 1. Who does intake — 20 claim handlers process claims; 10 support staff besides. (GGW (P. Gossmann), 8 Sep 2026)
- 2. Failed EASY imports — Not noticed today — so detection comes before retries. (GGW (P. Gossmann), 8 Sep 2026)
- 3. Volumes and shares — ~50 first reports a day · 65% minor · ~30 in the team. (GGW case study brief)
- 4. Wage base — Schadenregulierer/in, €45,100 gross — the role doing intake. (StepStone)
- 5. Employer on-cost — ≈ €23 per €100 gross — health, pension, care. (sevdesk)
- 6. Working time & breaks — 38h tariff week; statutory break rules. (Arbeitszeitgesetz §4)
- 7. Build & upkeep pay — Senior ML engineer ≈ €98k; juniors ≈ €70k, as briefed. (Glassdoor)
- 8. Model prices — gpt-5.6 terra — \$2 in / \$12 out per million tokens. (OpenAI)
- 9. Hosting & database — Container Apps consumption rates; PostgreSQL B2s. (Azure pricing)
- 10. Upkeep after go-live — 15–20% of build effort a year, ours is 17%. (Standard maintenance rule of thumb)
- 11. Industry baseline — 17 of 20 carriers still manual; none let AI decide alone. (Decerto, 2026 Claims Decisioning Pulse)
- 12. Where value comes from — The 10-20-70 split — algorithms, technology, people. (BCG)
- 13. Fixing import failures — Retry vs. dead-letter, idempotency, alert on rate. (Standard integration pattern)

Still assumptions, not sourced. The first thing we would test:

Every per-step minute on slide 8 · the 30% allowance for breaks and interruptions · how often a CPR is missing (30%) or underwriting is asked (25%) · how often the redesign raises an exception (20%) · the tokens a typical claim needs · the €146k build · the 0.25 FTE upkeep. Two weeks of watching live intake would replace all of this with measurements.

Prices checked September 2026 · US\$ converted at €0.92 · every euro figure is fully loaded (gross pay plus employer contributions).