

Redesigning intake gives each of the 20 claim handlers back 52 minutes a day

17.5 hours a day across the 20 of them — 2.8 full-time people · two thirds of that work goes to claims that were always going to close quietly · a person still decides anything adverse

TODAY

What one claim costs one claim handler

40 minutes per claim

- Re-reads the report, opens attachments
- Chases a missing CPR
- Switches EASY · IDB · e-mail · paper
- +30% for breaks and interruptions

REDESIGNED

The same checks and the same exits, reordered

19 minutes per claim

- 52% cut, not the 75% max
- 4 AI agents, 1 rules-only step
- 2 human checkpoints kept
- Handlers re-check before they trust it

WHERE THE EFFORT ACTUALLY GOES

- 5%** **Models and machines**
€0.08 a claim · €2,400 a year
- 25%** **Plumbing and audit trail**
EASY/IDB wiring, logging every decision
- 70%** **Process, training, trust**
adoption is the real schedule risk

Follows BCG's 10-20-70 rule, adjusted for this process

52 min

freed per claim handler, every working day (of 6.4 productive hours)

17.5 h

across the 20 claim handlers — 2.8 full-time people

€153k

value of that capacity, per year

14 mths

to pay back the one-off build of €146k

“€153k” is capacity, not cash: 4,370 handler-hours a year at the loaded rate. It becomes money once those hours go to the backlog.

So what: cost isn't the decision — value covers it several times over. Governance is: a person reviews every call that can go against a claimant.


A conservative estimate throughout — not the best case. Build-up and sources: slides 8–9.

Three ways in and a separate side-track diagram become one classified intake

FNOL as it runs today (left) against the redesign (right) · first report to an opened, policy-linked claim


Trigger / start Step / decision Exception / rework Exit

TODAY



AI agent (LLM call) Rules only, no model Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




These two agents together replace 3 entry channels, the auto-import failure path, and the entire 1.1 side-track diagram.

Three checks done one after another become one check done all at once

Is it complete, linked to the right policy, and not already in the system? None of the three depends on the others


- Trigger / start
- Step / decision
- Exception / rework
- Exit

TODAY



- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




The AI writes the draft; the handler still makes the decision

From a system warning to a closed claim or a fully opened one — nothing that can go against a claimant is automated

- Trigger / start
- Step / decision
- Exception / rework
- Exit


TODAY



Today: 7 decision gates, 3 entry channels, 10 exception or exit branches across the full FNOL flow.

- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail



Net result: 3 entry channels → 1 · 3 sequential gates → 1 parallel check · the Coverage Check handoff disappears, it runs on the same state

Start where a mistake is cheap — that is what earns the trust the bigger wins depend on

“Hard” means build effort and regulatory exposure together, not just weeks of work




- THE IMPORT FAILURES**
- Not noticed today (GGW) — so detection comes before retries:*
- 1 Count what should arrive**
daily reconciliation catches gaps
 - 2 Sort the failure**
timeout vs. missing field
 - 3 Retry with backoff**
1s, 2s, 4s + jitter
 - 4 One key per claim**
no duplicate replays
 - 5 Park, don't drop**
holding queue for review
 - 6 Alert on the rate**
20/hour ≠ 1/day

AES rulings and underwriting authority are deliberately absent — those decisions are not ours to make · frequencies from the case study


System at a glance — four lanes, seven stations, one shared claim state

SYSTEM AT A GLANCE



AI owns the drafting. Rules own the facts. A person owns every call that can go against a claimant.

One shared claim record carries every field and decision through the whole run, so a failed step resumes rather than restarts.




THE STACK




Deployed to the cloud GGW already use, with a switch to turn any agent off

OPERATING MODEL



HOW IT GOES LIVE



If an agent fails or is unsure, the claim simply falls back to today's manual queue — slower, never wrong.

Two numbers carry this case: the minutes saved, and what they're worth

A · TIME PER CLAIM, PER HANDLER (MINUTES)

Step	Today	Redesigned
Read & extract the report	16.1	3.5
Check policy, CPR & escalations	6.0	0.6
Decide, draft and send	5.5	3.0
Switch between systems	3.2	0.7
Touch time	30.8	7.8
+30% breaks & interruptions	9.2	2.3
Standard time	40.0	10.1
What we bank	40.0	19.0

We bank 19 min, not the 10.1 the design allows — a conservative three-point estimate, not the best case.

Basis: 20 claim handlers × 2.5 claims a day, confirmed by GGW (8 Sep 2026) — 52 min back, each, per day.

B · THE BUSINESS CASE

EVERY YEAR	
Capacity released	4,370 handler-h · 2.8 FTE
Worth, at €34.92 per handler-hour	€152,600
Cost to run it — AI models and Azure	– €2,400
Ongoing support — 0.25 of an engineer	– €24,400
Net value, every year	€125,800
ONE-OFF	
Build — 3 people × 6 months, fully loaded	€146,400
Payback	14 months

Running cost in full: ≈ €0.08/claim in AI tokens + ≈ \$134/month in Azure → €2,400/year, all in.

The build is loaded team capacity, not new cash — the AI team is already funded.

- 1. Who does intake — 20 claim handlers process claims; 10 support staff besides. (GGW (P. Gossmann), 8 Sep 2026)
- 2. Failed EASY imports — Not noticed today — so detection comes before retries. (GGW (P. Gossmann), 8 Sep 2026)
- 3. Volumes and shares — ~50 first reports a day · 65% minor · ~30 in the team. (GGW case study brief)
- 4. Wage base — Schadenregulierer/in, €45,100 gross — the role doing intake. (StepStone)
- 5. Employer on-cost — ≈ €23 per €100 gross — health, pension, care. (sevdesk)
- 6. Working time & breaks — 38h tariff week; statutory break rules. (Arbeitszeitgesetz §4)
- 7. Build & upkeep pay — Senior ML engineer ≈ €98k; juniors ≈ €70k, as briefed. (Glassdoor)
- 8. Model prices — gpt-5.6 terra — \$2 in / \$12 out per million tokens. (OpenAI)
- 9. Hosting & database — Container Apps consumption rates; PostgreSQL B2s. (Azure pricing)
- 10. Upkeep after go-live — 15–20% of build effort a year — ours is 17%. (Standard maintenance rule of thumb)
- 11. Industry baseline — 17 of 20 carriers still manual; none let AI decide alone. (Decerto, 2026 Claims Decisioning Pulse)
- 12. Where value comes from — The 10-20-70 split — algorithms, technology, people. (BCG)
- 13. Fixing import failures — Retry vs. dead-letter, idempotency, alert on rate. (Standard integration pattern)

Still assumptions — not sourced, and the first thing we would test:

Every per-step minute on slide 8 · the 30% allowance for breaks and interruptions · how often a CPR is missing (30%) or underwriting is asked (25%) · how often the redesign raises an exception (20%) · the tokens a typical claim needs · the €146k build · the 0.25 FTE upkeep. Two weeks of watching live intake would replace all of this with measurements.