


EXECUTIVE SUMMARY

Redesigning intake gives the 10-person support team back 1.75 hours each, every day

17.5 hours a day across the team — 2.8 full-time people · two thirds of that work goes to claims that were always going to close quietly · a person still decides anything adverse




Conservative throughout: the 21-minute-per-claim saving is a three-point estimate (best 30 / likely 21 / worst 12), not the design maximum. Build-up on slide 7, sources on slide 8.

Three ways in and a separate side-track diagram become one classified intake

FNOL as it runs today (left) against the redesign (right) · first report to an opened, policy-linked claim


Trigger / start Step / decision Exception / rework Exit

TODAY



AI agent (LLM call) Rules only, no model Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




These two agents together replace 3 entry channels, the auto-import failure path, and the entire 1.1 side-track diagram.

Three checks done one after another become one check done all at once

Is it complete, linked to the right policy, and not already in the system? None of the three depends on the others


- Trigger / start
- Step / decision
- Exception / rework
- Exit

TODAY



- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




The AI writes the draft; the handler still makes the decision

From a system warning to a closed claim or a fully opened one — nothing that can go against a claimant is automated

Trigger / start Step / decision Exception / rework Exit


TODAY



Today: 7 decision gates, 3 entry channels, 10 exception or exit branches across the full FNOL flow.

AI agent (LLM call) Rules only, no model Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




Net result across all three parts:


4 LLM agents + 1 rules-only node · 2 human checkpoints kept by design · 3 entry channels → 1 · 3 sequential gates → 1 parallel check
Coverage Check runs on the same state, beside claim creation — so the handoff between the two workflows disappears rather than being managed

Start where a mistake is cheap — that is what earns the trust the bigger wins depend on

“Hard” means build effort and regulatory exposure together, not just weeks of work



AES rulings and underwriting authority are deliberately absent — those decisions are not ours to make · frequencies from the case study




THE STACK



Deployed to the cloud GGW already use · infrastructure as code · one on/off switch per agent · model version pinned

HOW IT GOES LIVE



If an agent fails or is unsure, the claim drops into today's manual queue — slower, never wrong · what it costs is on slide 7

A · INTAKE TIME — MINUTES PER CLAIM, PER PERSON

Step	Today	Redesigned
Read the report, re-read it, open attachments	10.0	0.7
Take notes, check fields, chase a missing CPR	6.1	2.8
Attach policy · duplicate check · coverage date	5.0	0.0
Decide, draft and send / open the full claim	5.5	3.0
Work an escalation or an underwriting hold	1.0	0.6
Switch between EASY, IDB, e-mail and paper	3.2	0.7
Touch time	30.8	7.8
Breaks, fatigue, interruptions, rework (+30%)	9.2	2.3
Standard time — what the design allows	40.0	10.1
What the business case actually banks	40.0	19.0

The design allows 10.1 min a claim. We do not bank that. The case uses a three-point estimate — (best 30 + 4 × likely 21 + worst 12) ÷ 6 = 21.0 min saved a claim — which is the 19.0 min above, and the 19 minutes on slide 1. In plain terms: we assume handlers capture 70% of the saving the design allows.

Cross-check: 40 min × 12,500 claims a year = 8,332 hours — 52% of the 10-person support team's year, 17% of all 30. Per person: 5 claims a day at 40 min is 3.3 of a 6.4-hour day. Intake is half of what the support team does.

D · THE BUSINESS CASE

EVERY YEAR

Capacity released4,370 support-h · 2.8 FTE

Worth, at €31.83 per support-team hour€139,100 a year

Cost to run it — AI models and Azure– €2,400

Ongoing support — 0.25 of an engineer– €24,400

Net value, every year€112,200

ONE-OFF

Build — 3 people × 6 months, fully loaded€146,400

Payback16 months

The build is loaded team capacity, not new cash — the AI team is already funded. New cash needed: €2,400 a year. A shorter build costs less: 4 months ≈ €98k.

B · AI COST TO PROCESS ONE CLAIM

Agent	Tokens in	Out	Runs	US\$ / claim
Extraction	11,500	700	100%	0.0314
Classification	900	120	100%	0.0032
Validation — rules only, no model	—	—	100%	0.0000
Exception research	2,500	400	20%	0.0020
Outcome drafting	2,200	550	100%	0.0110
Subtotal				0.0476
× 1.75 for retries, guardrails, evals				0.0833
Per claim				≈ €0.077

One claim ≈ a scanned form, an e-mail and 3–4 attachments. All four agents run on one model tier — gpt-5.6 terra, \$2 in / \$12 out per million tokens.

C · AZURE, WEST EUROPE — MONTHLY

Component	US\$/mth
4 container apps — 2 always on, 2 burst only	55
PostgreSQL B2s + 64 GiB storage and backup	61
Blob storage for documents + audit logs	8
Net new, incl. ~8% West Europe premium	134

Registry, networking, Key Vault and DevOps (≈ €670/yr) already exist, so they are not charged here. Azure DevOps gives repos and pipelines, not a database — PostgreSQL is genuinely new.

E · WHO GETS THE TIME BACK — CONFIRMED BY GGW
20 claim handlers and 10 support staff (GGW, 8 Sep). The brief puts FNOL with the support team, so every freed hour is theirs.

Per support-team member	Today	Redesigned	Difference
Claims taken in, per day	5	5	—
Time on intake, per day	3.3 h	1.6 h	– 1.75 h
Share of a productive day	52%	25%	– 27 pts
Hours back, each, per day			1.75 h

Who does intake

20 claim handlers, 10 support staff — FNOL sits with support

GGW (P. Gossmann), 8 Sep 2026

Failed EASY imports

Not noticed today — so detection comes before retries

GGW (P. Gossmann), 8 Sep 2026

Volumes and shares

~50 first reports a day · 65% minor · ~30 in the team

GGW case study brief

Wage base — intake

Sachbearbeiter/in Versicherung, €41,100 gross

[StepStone Gehaltsreport](#)

Employer on-cost

≈ €23 per €100 gross — health, pension, care

[Lohnnebenkosten 2026, sevdesk](#)

Working time, breaks

38 h tariff week; statutory break rules

[Arbeitszeitgesetz §4](#)

Build & upkeep pay

Senior ML engineer ≈ €98k; juniors at €70k as briefed

[Glassdoor, ML Engineer Germany](#)

Model prices

gpt-5.6 terra — \$2 in / \$12 out per million tokens

[OpenAI API pricing](#)

Container hosting

Consumption plan rates and the monthly free grant

[Azure Container Apps pricing](#)

Database

PostgreSQL Flexible Server B2s — 2 vCore / 4 GiB

[Azure Database for PostgreSQL](#)

Upkeep after go-live

15–20% of build effort a year — our 0.25 FTE is 17%

Software maintenance rule of thumb

Industry baseline

17 of 20 carriers still manual; none let AI decide alone

[Decerto, 2026 Claims Decisioning Pulse](#)

Where value comes from

The 10-20-70 split — algorithms, technology, people

[BCG, AI in customer service operations](#)

Fixing import failures

Retry vs. dead-letter, idempotency keys, alert on rate

[Dead-letter queues and poison messages](#)

STILL OUR ASSUMPTIONS — AND THE FIRST THING WE WOULD TEST

Every per-step minute on slide 7 · the 30% allowance for breaks and interruptions · how often a CPR is missing (30%) or underwriting is asked (25%) · how often the redesign raises an exception (20%) · the tokens a typical claim needs · the €146k build · the 0.25 FTE upkeep.

Who does intake is no longer among them — GGW answered it, and every per-person figure in this deck is rebuilt on that answer. Two weeks of watching live intake would replace the rest with measurements.

confirmed by GGW or stated in the brief

[published source, linked](#)

Prices checked September 2026 · US\$ converted at €0.92 · every euro figure is fully loaded — gross pay plus employer contributions