

Redesigning intake gives each of the 20 claim handlers back 52 minutes a day

17.5 hours a day across the 20 of them, or 2.8 full-time people. Two thirds of that work goes to claims that were always going to close quietly.

TODAY

What one claim costs one claim handler

40 minutes per claim

- Re-reads the report, opens attachments
- Chases a missing CPR
- Switches EASY · IDB · e-mail · paper
- +30% for breaks and interruptions

REDESIGNED

The same checks and the same exits, reordered

19 minutes per claim

- 52% cut, not the 75% max
- 4 AI agents, 1 rules-only step
- 2 human checkpoints kept
- Handlers re-check before they trust it

WHERE THE EFFORT ACTUALLY GOES

- 5%** **Models and machines**
€0.08 a claim · €2,400 a year
- 25%** **Plumbing and audit trail**
EASY/IDB wiring, logging every decision
- 70%** **Process, training, trust**
adoption is the real schedule risk

Follows BCG's 10-20-70 rule, adjusted for this process

52 min

freed per claim handler, every working day (of 6.4 productive hours)

17.5 h

across the 20 claim handlers, or 2.8 full-time people

€153k

value of that capacity, per year

14 mths


to pay back the one-off build of €146k

“€153k” is capacity, not cash: 4,370 handler-hours a year at the loaded rate. It becomes money once those hours go to the backlog.

So what: cost isn't the decision. Value covers it several times over. Governance is: a person reviews every call that can go against a claimant.

Seven stations share one claim record, so nothing is handed off or read twice

SYSTEM AT A GLANCE




AI drafts. Rules check. A person decides anything that can go against a claimant.

Three ways in and a separate side-track diagram become one classified intake

Today on the left, the redesign on the right


- Trigger / start
- Step / decision
- Exception / rework
- Exit

TODAY



- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE one shared claim record, one audit trail




These two agents replace 3 entry channels, the auto-import failure path, and the whole 1.1 side track.


Three checks done one after another become one check done all at once

None of the three checks depends on the other two

TODAY




AI-NATIVE one shared claim record, one audit trail



The AI writes the draft; the handler still makes the decision


Nothing that can go against a claimant is automated

TODAY



Today: 7 decision gates, 3 entry channels, 10 exception or exit branches across the full FNOL flow.


AI-NATIVE *one shared claim record, one audit trail*




Net result: 3 entry channels → 1 · 3 sequential gates → 1 parallel check · the Coverage Check handoff disappears, it runs on the same state

Start where a mistake is cheap. That is what earns the trust the bigger wins need.

“Hard” means build effort and regulatory exposure together, not just weeks of work




AES rulings and underwriting authority are deliberately absent. Those decisions are not ours to make.




THE STACK



OPERATING MODEL



HOW IT GOES LIVE



If an agent fails or is unsure, the claim falls back to today's manual queue. Slower, never wrong.

Every minute is observed, and the redesign pays for itself in 14 months

A · TIME PER CLAIM, PER HANDLER (MINUTES)

Step	Today	Redesigned
Read & extract the report	16.1	3.5
Check policy, CPR & escalations	6.0	0.6
Decide, draft and send	5.5	3.0
Switch between systems	3.2	0.7
Touch time	30.8	7.8
+30% breaks & interruptions	9.2	2.3
Standard time	40.0	10.1
What we bank	40.0	19.0

We bank 19 min, not the 10.1 the design allows. Conservative three-point estimate, not the best case.

Basis: 20 claim handlers × 2.5 claims a day, confirmed by GGW (8 Sep 2026). 52 min back, each, per day.

B · THE BUSINESS CASE

EVERY YEAR	
Capacity released	4,370 handler-h · 2.8 FTE
Worth, at €34.92 per handler-hour	€152,600
Cost to run it — AI models and Azure	– €2,400
Ongoing support — 0.25 of an engineer	– €24,400
Net value, every year	€125,800
ONE-OFF	
Build — 3 people × 6 months, fully loaded	€146,400
Payback	14 months

Running cost in full: ≈ €0.08/claim in AI tokens + ≈ \$134/month in Azure → €2,400/year, all in.

The build is loaded team capacity, not new cash. The AI team is already funded.

Every number is sourced, dated, or named as an assumption we'd test first

- 1. Who does intake — 20 claim handlers process claims; 10 support staff besides. (GGW (P. Gossmann), 8 Sep 2026)
- 2. Failed EASY imports — Not noticed today — so detection comes before retries. (GGW (P. Gossmann), 8 Sep 2026)
- 3. Volumes and shares — ~50 first reports a day · 65% minor · ~30 in the team. (GGW case study brief)
- 4. Wage base — Schadenregulierer/in, €45,100 gross — the role doing intake. (StepStone)
- 5. Employer on-cost — ≈ €23 per €100 gross — health, pension, care. (sevdesk)
- 6. Working time & breaks — 38h tariff week; statutory break rules. (Arbeitszeitgesetz §4)
- 7. Build & upkeep pay — Senior ML engineer ≈ €98k; juniors ≈ €70k, as briefed. (Glassdoor)
- 8. Model prices — gpt-5.6 terra — \$2 in / \$12 out per million tokens. (OpenAI)
- 9. Hosting & database — Container Apps consumption rates; PostgreSQL B2s. (Azure pricing)
- 10. Upkeep after go-live — 15–20% of build effort a year, ours is 17%. (Standard maintenance rule of thumb)
- 11. Industry baseline — 17 of 20 carriers still manual; none let AI decide alone. (Decerto, 2026 Claims Decisioning Pulse)
- 12. Where value comes from — The 10-20-70 split — algorithms, technology, people. (BCG)
- 13. Fixing import failures — Retry vs. dead-letter, idempotency, alert on rate. (Standard integration pattern)

Still assumptions, not sourced. The first thing we would test:

Every per-step minute on slide 8 · the 30% allowance for breaks and interruptions · how often a CPR is missing (30%) or underwriting is asked (25%) · how often the redesign raises an exception (20%) · the tokens a typical claim needs · the €146k build · the 0.25 FTE upkeep. Two weeks of watching live intake would replace all of this with measurements.