

Redesigning intake releases 17.5 handler-hours a day — the capacity of 2.8 full-time people

Two thirds of that work goes to claims that were always going to close quietly · a person still decides anything that can go against a claimant

TODAY

What one claim costs one handler today, start to finish

40 minutes per claim

Reading and re-reading the report, opening attachments, chasing a missing CPR, switching between EASY, IDB, e-mail and paper. Plus 30% for breaks and interruptions.

REDESIGNED

The same checks and the same exits — reordered, some reassigned

19 minutes per claim

A 52% cut, not the 75% the design allows: handlers will re-check drafts before they trust them. Four AI agents, one rules-only step, two human checkpoints.

WHERE THE EFFORT ACTUALLY GOES

- 5%** **Models and machines**
€0.077 a claim
- 25%** **Plumbing and audit trail**
EASY/IDB wiring, logging every decision
- 70%** **Process, training, trust**
adoption is the real schedule risk

Follows BCG's 10-20-70 rule, adjusted for this process

17.5 h

handler-hours freed every working day, across the claims team

2.8 FTE

the same capacity, expressed as full-time people

€153k

value of that capacity, per year

14 mths

to pay back the one-off build of €146k

What “€153k” means: 4,370 handler-hours a year × €34.92, the real cost of one handler-hour. This is capacity, not cash. It becomes money only if the hours go to the backlog, or are absorbed by normal staff turnover.

So what: cost is not what decides this. The AI costs €2,400 a year to run — under 2% of the value it releases. What decides it is governance: proving a person reviewed anything that can go against a claimant, and whether handlers trust the drafts enough to approve them.


Conservative throughout: the 21.0-minutes-per-claim saving is a three-point estimate (best 30 / likely 21 / worst 12), not the design maximum. Build-up on slide 6, sources on slide 7.

Three ways in and a separate side-track diagram become one classified intake

FNOL as it runs today (left) against the redesign (right) · first report to an opened, policy-linked claim


Trigger / start Step / decision Exception / rework Exit

TODAY



AI agent (LLM call) Rules only, no model Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




These two agents together replace 3 entry channels, the auto-import failure path, and the entire 1.1 side-track diagram.

Three checks done one after another become one check done all at once

Is it complete, linked to the right policy, and not already in the system? None of the three depends on the others


- Trigger / start
- Step / decision
- Exception / rework
- Exit

TODAY



- AI agent (LLM call)
- Rules only, no model
- Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail



The AI writes the draft; the handler still makes the decision

From a system warning to a closed claim or a fully opened one — nothing that can go against a claimant is automated


- Trigger / start

Step / decision

Exception / rework

Exit

TODAY




Today: 7 decision gates, 3 entry channels, 10 exception or exit branches across the full FNOL flow.

- AI agent (LLM call)

Rules only, no model

Human checkpoint

AI-NATIVE — one LangGraph state graph · shared claim state carries the audit trail




Net result across all three parts:

4 LLM agents + 1 rules-only node · 2 human checkpoints kept by design · 3 entry channels → 1 · 3 sequential gates → 1 parallel check · shared claim state carries the audit trail

Start where a mistake is cheap — that is what earns the trust the bigger wins depend on

“Hard” means build effort and regulatory exposure together, not just weeks of work



AES rulings and underwriting authority are deliberately absent — those decisions are not ours to make · frequencies from the case study

A · HANDLER TIME — MINUTES PER CLAIM, PER HANDLER

Step	Today	Redesigned
Read the report, re-read it, open attachments	10.0	0.7
Take notes, check fields, chase a missing CPR	6.1	2.8
Attach policy · duplicate check · coverage date	5.0	0.0
Decide, draft and send / open the full claim	5.5	3.0
Work an escalation or an underwriting hold	1.0	0.6
Switch between EASY, IDB, e-mail and paper	3.2	0.7
Touch time	30.8	7.8
Breaks, fatigue, interruptions, rework (+30%)	9.2	2.3
Standard time — what the design allows	40.0	10.1
What the business case actually banks	40.0	19.0

The design allows 10.1 min a claim. We do not bank that. The case uses a three-point estimate — (best 30 + 4 × likely 21 + worst 12) ÷ 6 = 21.0 min saved a claim — which is the 19.0 min above, and the 19 minutes on slide 1.

Cross-check: 40 min × 12,500 claims a year = 8,332 handler-hours — 17% of a 30-person team's yearly capacity. Plausible for the front end of a claims process. Had it come out at 60%, the model would be wrong.

D · THE BUSINESS CASE

EVERY YEAR

Capacity released4,370 handler-h · 2.8 FTE

Worth, at €34.92 per handler-hour€152,600 a year

Cost to run it — AI models and Azure– €2,400

Ongoing support — 0.25 of an engineer– €24,400

Net value, every year€125,800

ONE-OFF

Build — 3 people × 6 months, fully loaded€146,400

Payback14 months

The build is loaded team capacity, not new cash — the AI team is already funded. New cash needed: €2,400 a year. A shorter build costs less: 4 months ≈ €98k.

B · AI COST TO PROCESS ONE CLAIM

Agent	Tokens in	Out	Runs	US\$ / claim
Extraction	11,500	700	100%	0.0314
Classification	900	120	100%	0.0032
Validation — rules only, no model	—	—	100%	0.0000
Exception research	2,500	400	20%	0.0020
Outcome drafting	2,200	550	100%	0.0110
Subtotal				0.0476
× 1.75 for retries, guardrails, evals				0.0833
Per claim				≈ €0.077

One claim ≈ a scanned form, an e-mail and 3–4 attachments. All four agents run on one model tier — gpt-5.6 terra, \$2 in / \$12 out per million tokens.

C · AZURE, WEST EUROPE — MONTHLY

Component	US\$/mth
4 container apps — 2 always on, 2 burst only	55
PostgreSQL B2s + 64 GiB storage and backup	61
Blob storage for documents + audit logs	8
Net new, incl. ~8% West Europe premium	134

Registry, networking, Key Vault and DevOps (≈ €670/yr) already exist, so they are not charged here. Azure DevOps gives repos and pipelines, not a database — PostgreSQL is genuinely new.

E · WHO ACTUALLY GETS THE TIME BACK

The brief says a support team does intake, not all 30 handlers — so the truth is in the middle rows. Our first question for GGW.

Intake handled by...	Claims each, per day	Intake work, per day	Time freed, per day
all 30 claims handlers	1.7	1.1 h	35 min
a support team of 12	4.2	2.8 h	88 min
a support team of 8	6.2	4.2 h	2.2 h
a support team of 6	8.3	5.6 h	2.9 h

Wage base Sachbearbeiter Schadenregulierung, €45,100 gross	StepStone Gehaltsreport
Employer on-cost ≈ €23 per €100 gross — health, pension, unemployment, care	Lohnnebenkosten 2026, sevdesk
Working time, breaks 38 h tariff week; statutory break rules	Arbeitszeitgesetz §4
Model prices gpt-5.6 terra — \$2 in / \$12 out per million tokens	OpenAI API pricing
Container hosting Consumption plan rates and the monthly free grant	Azure Container Apps pricing
Database PostgreSQL Flexible Server B2s — 2 vCore / 4 GiB	Azure Database for PostgreSQL
Build-team pay Senior ML engineer ≈ €98k; juniors at €70k as briefed	Glassdoor, ML Engineer Germany
Upkeep after go-live 15–20% of build effort a year — our 0.25 FTE is 17%	<i>Software maintenance rule of thumb</i>
Industry baseline 17 of 20 carriers still manual; none let AI decide alone	Decerto, 2026 Claims Decisioning Pulse
Where value comes from The 10-20-70 split — algorithms, technology, people	BCG, AI in customer service operations
Fixing import failures Retry vs. dead-letter, idempotency keys, alert on rate	Dead-letter queues and poison messages
Volumes and shares ~50 first reports a day · 65% minor · ~30 handlers	<i>GGW case study brief</i>

OUR ASSUMPTIONS — NOT SOURCED, AND THE FIRST THING WE WOULD TEST

Every per-step minute on slide 6 · the 30% allowance for breaks and interruptions · how often a CPR is missing (30%) or underwriting is asked (25%) · how often the redesign raises an exception (20%) · the tokens a typical claim needs · the €146k build · the 0.25 FTE upkeep · and how many people staff intake, which sets the per-person saving.

Two weeks of watching live intake would replace all of these with real measurements. That is step one of the plan, not an afterthought.

Prices checked September 2026 · US\$ converted at €0.92 · every euro figure is fully loaded — gross pay plus employer contributions